Research Statement

Xiang (Jenny) Ren

Performance makes or breaks a software system. Severe performance issues render systems unusable; Even small performance degradation can be costly— Google finds that a 0.5-second delay in page load time caused a 20% drop in repeat traffic [14]. Despite its importance, there are many obstacles to building performant systems. Modern software systems have complex performance characteristics that are difficult to understand, and users and developers still invest significant time manually diagnosing critical performance issues.

The goal of my research is to help developers build performant software systems. My approach is to understand the system performance characteristics comprehensively, then examine where they break existing assumptions and identify opportunities to create automated techniques to improve performance.

For my dissertation, I carried out a study [15] that discovered many of Linux's core kernel operations had slowed down over time or had unstable performance, due to adding new features and security guarantees, as well as misconfigurations. Despite unavoidable performance trade-offs, I found much of the slowdown can be improved through proactive testing, diagnosis and optimization, or avoided through custom kernel configuration. The study motivated me to create an automated tool to diagnose the root causes of performance issues. I recognized that existing root causes diagnosis tools make fundamental assumptions that are broken by performance issues. Such tools only work when the bug has clear-cut execution outcomes *e.g.*, crashes or no crashes, and *absolute* root causes, *e.g.*, broken invariants or predicates; when the expected performance outcomes are *relative* to inputs workloads and performance root causes tend to be abnormal *relative* occurrences of events. I designed relational debugging that caters to the "relative" property of performance issues and can effectively diagnose complex performance issues that human developers struggle with. In addition to my main thesis work, I have also worked on improving system reliability, and on building performance systems for persistent memory.

For future work, I am inspired by the observations from my Linux performance study to create an automated pipeline for performance bug detection, diagnosis, and patching. Given that existing tools that serve similar purposes are often built for functional failures, I see potential for novel techniques that cater to the characteristics of performance issues. In addition to creating an external tool chain, I am also interested to identify ways that software systems can be modified internally such that they are easier to observe for the purpose of performance debugging. Finally, I am excited to identify new opportunities to improve performance by uncovering unknown performance phenomena and by studying emerging workloads.My current project is on studying performance creep, which is challenging to diagnose because it is a form of gradual degradation but results in significant impact in the long run.

A study of the performance evolution of Linux's Core Operations. I carried out a study to understand the performance of Linux's core operations, such as system calls and context switches, as the system evolved for seven and half years [15]. Surprisingly, the performance of many core operations had worsened or fluctuated significantly over time. I found the increase in overhead was caused by the need to accommodate an increasing number of new features and security enhancements. Further, simple misconfigurations had also caused significant performance fluctuations. I attributed the causes of these fluctuations and slowdowns to 11 changes. By disabling these root causes, I was able to speed up Redis, Apache, and Nginx benchmark workloads by as much as 56%, 33%, and 34%, respectively.

Moreover, I identified performance costs paid in addition to those required to increase security and accommodate new features. Out of the 11 root causes, four can be optimized and five can be avoided through better configuration. I found optimizations for two root causes and better configurations for two more, In addition to those eventually carried out Linux/Ubuntu developers, These findings suggest that better performance can be achieved through proactive testing, diagnosis, and optimization, and custom configuring kernels can help achieve a balance between performance, functionality and security.

Impact The findings from the study sparked many discussions on Hacker News and Reddit, and the paper was featured by the *Morning Paper* [5]. The benchmark designed for the study was deployed by Amazon.

Perspect: Automating the diagnosis of the root causes of performance issues. While carrying out the previous study, I manually diagnosed the performance changes and found no readily available tools that automate performance diagnosis. Profilers identify performance symptoms, such as hotspots, but do not link them to the root causes. Existing root cause diagnosis tools are *not* built for performance issues—they work for bugs that cause execution outcomes deemed incorrect regardless of inputs, *e.g.*, crashes; However, the expected performance outcome is *relative* to inputs workloads. In addition, existing tools capture *absolute* forms of root causes, like violated invariants or predicates, when performance root causes also tend to be *relative*, like an event happening more frequently.

Based on this observation, I invented relational debugging. Inspired by the idea of choosing the right frames of observation in the physical world, relational debugging chooses the appropriate fine-grained events as the basis of analysis, and captures performance root causes in terms of occurrences relative to these events. This design works independent of the effects of changing inputs and is generalizable to diverse forms of performance root causes.

We implemented relational debugging with Perspect [16]. Perspect is able to automatically and efficiently pinpoint the root causes of complex performance bugs. I used Perspect to diagnose Go-909 [1], which took developers a year to diagnose through a trial-and-error process.

Impact I used Perspect to diagnose two challenging open bugs (MongoDB-57221 [6], MongoDB-56274 [7]), and the root causes reported by Perspect were confirmed by MongoDB developers. A developer commented that "(Perspect's result) ties all the pieces together into a nice explanation." Perspect is requested by developers of the MMTk memory management framework and has attracted interest from CockroachDB.

Improving the reliability of distributed systems. I have also worked on projects aimed at improving the reliability of distributed systems, including Pensieve [18], which automatically constructs near-minimal reproductions of failures with the scalable event-chaining approach. During my internship at Microsoft Research, I applied the event-chaining algorithm to automatically reproduce flaky tests that are caused by concurrency [12]. The prototype is able to reproduce 26 out of 31 flaky tests in production codebases in under 6 minutes on average.

Building performant systems for emerging hardware. I have also worked on building novel file systems for persistent memory [13]. ctFS removes file indexing overhead and directly manages persistent memory using the hardware memory management unit (MMU) by allocating files as contiguous virtual memory chunks.

1 Future Research

For future work, I plan to take a multi-pronged approach to improving system performance. (1) I'm interested in studying the performance of emerging workloads, including serverless computing and machine learning workloads. I'm also interested in uncovering unknown performance phenomena in system software in general, including database software, distributed systems, operating systems, etc. (section 1.1). Through the studies, I plan to re-examine fundamental assumptions around system performance and identify opportunities for better system design or automated techniques that enhance performance testing, diagnosis, or optimization. (2) I plan to build more observable software systems to enhance performance diagnosability (section 1.2) and (3) create novel techniques that form a pipeline that automates performance bug detection, diagnosis, and fixes (section 1.3). My vision is to create a self-adaptive system that detects performance issues and creates patches or reconfigures the system automatically.

1.1 Uncovering and studying new performance phenomena

Understanding performance creep. I'm currently carrying out a comprehensive study of performance creep. Performance creep refers to gradual performance degradation that is allowed to persist over time; however, its impact is often significant in the long run. For example, A basic `SELECT` workload in MySql [10] experienced a 34% regression from version 4.1.22 to 5.7.2 [2]. We find creep also causes significant regression over time in popular systems, including MongoDB [9] and GCC [8]. Unfortunately, the causes of creep remain largely mysterious. The main reason is that performance creep is very challenging to diagnose, as creep consists of fine-grained changes, each of which is very short-running and easily confused with background noise from the OS, etc.; Existing profiling strategies are inadequate as they are too coarse-grained or have too much overhead. I plan to identify and develop techniques to measure small

performance differences accurately, including execution noise reduction and high-precision profiling. Subsequently, I plan to comprehensively study the root causes of creep and identify opportunities to optimize performance. Ultimately, my goal is to enable creep to be automatically detected, diagnosed, and optimized.

1.2 Enhancing performance diagnosability by building more observable systems

Capturing the root causes of performance outliers in situ, and across software components. Performance outliers are impactful yet elusive. They are hard to reproduce and capture using profilers [11]. Further, built-in checks, such as asserts, are also limited because they only work for *absolute* conditions; When debugging performance issues, whether the execution is expected must be determined *relative* to input workloads. Similar to relational debugging [16], I propose *relational asserts*, which check for the expected execution relative to fine-grained program events and hold regardless of input workloads. I propose to automatically learns what to assert relationally by observing normal executions.

Moreover, different software components often cannot agree on the same set of performance expectations. For example, Go-8832 [4] was famously hard to diagnose because the memory bloat was *not* caused by the Go application but by the OS, which excessively promoted hugepages. With relational asserts systematically inserted into an application, I propose to automatically translate these asserts into different software components. For Go-8832, relational asserts on memory usage in the Go application can be translated into the OS to check the frequency of hugepage promotions.

Embedding execution information into built-in performance metrics. Real-world software systems have large collections of built-in performance metrics to aid in debugging. For example, the go runtime has 70 performance metrics [3]. Developers often look for abnormal performance metrics when diagnosing performance issues, but performance metrics alone have limited explanatory power. For example, a developer may have observed an abnormal `heap_size` but has no information on the execution that led to it; Further diagnosis with very little lead can be challenging, like in the case of Go-909 [1]. I plan to design informative performance metrics that maximally encode relevant diagnostic information, such as variable values and execution paths while causing minimal execution overhead.

1.3 Building a pipeline of performance bug detection & diagnosis & patch generation

Detecting performance bugs before they cause severe consequences. My previous work [16] found that for many of the performance bugs, the root causes are also triggered in "healthy" runs, only less often. For example, the memory leak bug from Go-909 [1] also occurred in the good run. For future work, I plan to design an automated tool to detect such "hidden" root causes and identify input workloads that can cause them to manifest into performance issues.

Extending relational debugging to diverse performance diagnosis scenarios. Relational debugging is currently implemented with Perspect [16], which diagnoses the root causes of performance issues offline in a single-machine setting. I plan to build an online root cause diagnosis tool that carries out relational debugging incrementally to achieve low overhead. Secondly, I plan to adapt relational debugging to work on distributed systems as well as the Linux kernel.

Coupling performance diagnosis with patch generation. I find that well-known patterns of non-optimal performance—such as lack of caching, redundant work, sub-optimal configurations, etc. [17]—correlate with different types of root causes identified by Perspect [16]. For future work, I'm interested in creating a pipeline where the performance root causes diagnosed by Perspect are then used as feedback to automatically generate patches in a principled manner.

References

[1] 2010. runtime: garbage collection ineffective on 32-bit 909. (2010). `https://github.com/golang/go/issues/909`

[2] 2013. [MDEV-5333] Single-threaded performance regressions - Jira. (2013). `https://jira.mariadb.org/browse/MDEV-5333`

[3] 2013. [runtime/metrics. (2013). `https://pkg.go.dev/runtime/metrics`

[4] 2014. runtime: GC freeing goroutines memory but then taking it again 8832. (2014). `https://github.com/golang/go/issues/8832`

[5] 2019. An analysis of performance evolution of Linux's core operations. (2019). `https://blog.acolyer.org/2019/11/04/an-analysis-of-performance-evolution-of-linuxs-core-operations/`

[6] 2021. Inconsistent delete performance on collections with multiple secondary indexes on same key. (2021). `https://jira.mongodb.org/browse/SERVER-57221`

[7] 2021. TTL deletes are much slower on descending indexes than ascending indexes. (2021). `https://jira.mongodb.org/browse/SERVER-56274`

[8] 2023. GCC, the GNU Compiler Collection. (2023). `https://gcc.gnu.org/`

[9] 2023. MongoDB: The Developer Data Platform — MongoDB. (2023). `https://www.mongodb.com/`

[10] 2023. MySQL. (2023). `https://www.mysql.com/`

[11] Alexandra Fedorova, Craig Mustard, Ivan Beschastnikh, Julia Rubin, Augustine Wong, Svetozar Miucin, and Louis Ye. 2018. Performance Comprehension at WiredTiger. In *Proceedings of the 2018 26th ACM Joint Meeting on European Software Engineering Conference and Symposium on the Foundations of Software Engineering (ESEC/FSE 2018)*. Association for Computing Machinery, New York, NY, USA, 83–94. `https://doi.org/10.1145/3236024.3236081`

[12] Tanakorn Leesatapornwongsa, Xiang Ren, and Suman Nath. 2022. FlakeRepro: Automated and Efficient Reproduction of Concurrency-Related Flaky Tests. In *Proceedings of the 30th ACM Joint European Software Engineering Conference and Symposium on the Foundations of Software Engineering (ESEC/FSE 2022)*. Association for Computing Machinery, New York, NY, USA, 1509–1520. `https://doi.org/10.1145/3540250.3558956`

[13] Ruibin Li, Xiang Ren, Xu Zhao, Siwei He, Michael Stumm, and Ding Yuan. 2022. ctFS: Replacing File Indexing with Hardware Memory Translation through Contiguous File Allocation for Persistent Memory. In *20th USENIX Conference on File and Storage Technologies (FAST 22)*. USENIX Association, Santa Clara, CA, 35–50. `https://www.usenix.org/conference/fast22/presentation/li`

[14] Greg Linden. 2017. Marissa Mayer at Web 2.0. `http://glinden.blogspot.com/2006/11/marissa-mayer-at-web-20.html`. (Nov. 2017).

[15] Xiang (Jenny) Ren, Kirk Rodrigues, Luyuan Chen, Camilo Vega, Michael Stumm, and Ding Yuan. 2019. An Analysis of Performance Evolution of Linux's Core Operations. In *Proceedings of the 27th ACM Symposium on Operating Systems Principles (SOSP '19)*. Association for Computing Machinery, New York, NY, USA, 554–569. `https://doi.org/10.1145/3341301.3359640`

[16] Xiang (Jenny) Ren, Sitao Wang, Zhuqi Jin, David Lion, Adrian Chiu, Tianyin Xu, and Ding Yuan. 2023. Relational Debugging — Pinpointing Root Causes of Performance Problems. In *17th USENIX Symposium on Operating Systems Design and Implementation (OSDI 23)*. USENIX Association, Boston, MA, 65–80. `https://www.usenix.org/conference/osdi23/presentation/ren`

[17] Guoqing Xu, Nick Mitchell, Matthew Arnold, Atanas Rountev, and Gary Sevitsky. 2010. Software Bloat Analysis: Finding, Removing, and Preventing Performance Problems in Modern Large-Scale Object-Oriented Applications. In *Proceedings of the FSE/SDP Workshop on Future of Software Engineering Research (FoSER '10)*. Association for Computing Machinery, New York, NY, USA, 421–426. `https://doi.org/10.1145/1882362.1882448`

[18] Yongle Zhang, Serguei Makarov, Xiang Ren, David Lion, and Ding Yuan. 2017. Pensieve: Non-Intrusive Failure Reproduction for Distributed Systems Using the Event Chaining Approach. In *Proceedings of the 26th Symposium on Operating Systems Principles (SOSP '17)*. Association for Computing Machinery, New York, NY, USA, 19–33. `https://doi.org/10.1145/3132747.3132768`